


Assignment 4 – Solution

Machine Learning
MSc Business Analytics
Wolfram Wiesemann

1 Individual Assignment


1. *Plot the observations in a two-dimensional graph.*

The graph looks as follows:



2. *Perform K -means clustering with $K = 2$ using the Euclidean norm. Toss a coin 7 times to initialise the algorithm.*

First we assign randomly $C_1 = \{2, 6, 7\}$ (red) and $C_2 = \{1, 3, 4, 5\}$ (black):




We compute the centroids of the two classes as $\mathbf{c}_1 = (\frac{11}{3}, 2)$ and $\mathbf{c}_2 = (\frac{9}{4}, \frac{5}{2})$. We now recompute the distances:

Obs. i	x_{i1}	x_{i2}	$\text{dist}(\mathbf{x}_i, \mathbf{c}_1)$	$\text{dist}(\mathbf{x}_i, \mathbf{c}_2)$
1	1	4	2.84	1.95
2	1	3	2.84	1.95
3	1	2	4.01	2.79
4	5	1	1.33	2.79
5	2	3	4.33	3.5
6	6	2	3.07	4.03
7	4	1	2.02	3.05

After reassignment, the new clusters are $C_1 = \{4, 6, 7\}$ and $C_2 = \{1, 2, 3, 5\}$. The new centroids of the two clusters are $\mathbf{c}_1 = (5, \frac{4}{3})$ and $\mathbf{c}_2 = (\frac{5}{4}, 3)$.

Obs. i	x_{i1}	x_{i2}	$\text{dist}(\mathbf{x}_i, \mathbf{c}_1)$	$\text{dist}(\mathbf{x}_i, \mathbf{c}_2)$
1	1	4	4.01	2.01
2	1	3	4.01	2.01
3	1	2	5.42	2.01
4	5	1	0.67	3.88
5	2	3	5.55	3.09
6	6	2	2.85	4.85
7	4	1	1.66	4.06

The clusters are still $C_1 = \{4, 6, 7\}$ and $C_2 = \{1, 2, 3, 5\}$. The algorithm thus terminates with the following result:



3. Cluster the data using hierarchical clustering with complete linkage and the Euclidean norm. Draw the resulting dendrogram.

We calculate the following pairwise distances between the observations:

	{1}	{2}	{3}	{4}	{5}	{6}	{7}
{1}	0						
{2}	1	0					
{3}	2	1	0				
{4}	5	4.47	4.12	0			
{5}	1.41	1	1.41	3.6	0		
{6}	5.38	5.09	5	1.41	4.12	0	
{7}	4.24	3.6	3.16	1	2.82	2.23	0

(Empty cells can be inferred from symmetry.) We first merge the ‘clusters’ {1} and {2}:

	{1, 2}	{3}	{4}	{5}	{6}	{7}
{1, 2}	0					
{3}	2	0				
{4}	5	4.12	0			
{5}	1.41	1.41	3.6	0		
{6}	5.38	5	1.41	4.12	0	
{7}	4.24	3.16	1	2.82	2.23	0

We now merge the ‘clusters’ {4} and {7}:

	{1, 2}	{3}	{5}	{6}	{4, 7}
{1, 2}	0				
{3}	2	0			
{5}	1.41	1.41	0		
{6}	5.38	5	4.12	0	
{4, 7}	5	4.12	3.6	2.23	0

We now merge the ‘clusters’ {5} and {1, 2}:

	{1, 2, 5}	{3}	{6}	{4, 7}
{1, 2, 5}	0			
{3}	2	0		
{6}	5.38	5	0	
{4, 7}	5	4.12	2.23	0


We now merge the ‘clusters’ {3} and {1, 2, 5}:

	{1, 2, 3, 5}	{6}	{4, 7}
{1, 2, 3, 5}	0		
{6}	5.38	0	
{4, 7}	5	2.23	0

We now merge the ‘clusters’ {6} and {4, 7}:

	{1, 2, 3, 5}	{4, 6, 7}
{1, 2, 3, 5}	0	
{4, 6, 7}	5.38	0

After merging the clusters $\{1, 2, 3, 5\}$ and $\{4, 6, 7\}$, we obtain the following result:



2 Group Assignment

1. *Explore manually the website <http://sofifa.com>. Under the tab ‘All players’, press on the Argentinian flag. Notice how the URL of the opened webpage changes to <http://sofifa.com/players?na=52>. Scrolling down, notice that not all players fit in one page. If you press ‘Next’, the new URL is <http://sofifa.com/players?na=52&offset=100>. Can you see the pattern? Next select an individual player and notice how the URL changes. We want to download the numerical attributes available for all 300 Argentinian players.*

By looking at a few pages, we see that $\text{na}=\mathbf{x}$ in the URL refers to the nationality, and that Argentina corresponds to $\mathbf{x}=52$. Using $\text{offset}=\mathbf{x}$, we can retrieve all players with numbers $\mathbf{x}, \dots, \mathbf{x} + 100$.

2. *Explain in detail the code below. In order to better understand the code, you may want to look at the following websites:*
 - <https://www.crummy.com/software/BeautifulSoup/>
 - <http://www.aivosto.com/vbtips/regex.html>
 - <https://docs.python.org/2/library/re.html>

We go through the code step by step:

```
import pandas as pd
from bs4 import BeautifulSoup
import requests
import re
import unicodedata
```

These commands import the employed libraries.

```
attributes = [ 'Crossing', 'Finishing', 'Heading_Accuracy',
               'Short_Passing', 'Volleys', 'Dribbling', 'Curve',
               'Free_Kick_Accuracy', 'Long_Passing', 'Ball_Control', 'Acceleration',
               'Sprint_Speed', 'Agility', 'Reactions', 'Balance',
               'Shot_Power', 'Jumping', 'Stamina', 'Strength',
               'Long_Shots', 'Aggression', 'Interceptions', 'Positioning',
               'Vision', 'Penalties', 'Composure', 'Marking',
               'Standing_Tackle', 'Sliding_Tackle', 'GK_Diving',
               'GK_Handling', 'GK_Kicking', 'GK_Positioning', 'GK_Reflexes' ]
```

These commands define a list that contains all attributes of interest.

```
links = [] # Download data for all 300 Argentinian players
for offset in [ '0', '100', '200' ]:
    page = requests.get ( 'http://sofifa.com/players?na=52&offset=' + offset )
    soup = BeautifulSoup (page.content, 'html.parser')
    for link in soup.find_all ( 'a' ):
        links.append (link.get ( 'href' ))
links = [ 'http://sofifa.com' + l for l in links if 'player/' in l ]
```

Here we loop over 3 pages as we want to retrieve 300 players, that is, 100 players per page. We use the `requests` package to obtain the HTML pages. The library `BeautifulSoup` is useful for extracting from each page all the player URLs. You can use

```
print soup.prettify()
```

to explore the object. The code above extracts all the URLs found within the `<a>` tags of an HTML page. There are some URLs that do not point to players: We filter them in the last line, where we also make the URLs absolute by adding `'http://sofifa.com'`.

```
# pattern for regular expression
pattern = r"""\s*([\w\s]*)""" # file starts with empty spaces... players
                             # name... - other stuff
for attr in attributes:
    pattern += r"""\d*\s*""" + attr + r"""""" # for each attribute we have
                                             # other stuff... number... attribute... other stuff
pat = re.compile (pattern, re.DOTALL) # parsing multiline text
```

This code block constructs a regular expression pattern that we need when looping over the players' webpages. We will further comment on this below.

```
rows = []
for j, link in enumerate (links):
    print j, link
    row = [link]
```

We loop over the accumulated links, one for every player. The purpose of the loop is to store the link, the player's name and all the attribute values for each player. We start by storing the link that is already available.

```
playerpage = requests.get (link)
playersoup = BeautifulSoup (playerpage.content, 'html.parser')
text = playersoup.get_text()
text = unicodedata.normalize ('NFKD', text).encode ('ascii', 'ignore')
```

Here we retrieve the HTML page for the player. The tag information is not very helpful, therefore we convert it to normal text and we are going to parse it using regular expressions. The last line replaces accented Latin characters that appear in some players' names with English letters. Such tedious steps are often necessary when scraping websites, emails and other forms of unstructured data. Getting things to work inevitably require a trial-and-error process. Running the commands from an interpreter for a specific player (link), you can explore the text with

```
print text
```

The output is of the following form:

```
Lionel Messi – FIFA 17 – Feb 14, 2017 – SoFIFA
.....
.....
77 Crossing

95 Finishing

71 Heading Accuracy
.....
```

The output is fairly consistent for all players. We note that the text starts with the player's name, followed by a '-'. At some point later in the file we get the pattern of a number followed by the attributes of interest. The pattern we have constructed before the loop captures the required fields:

```
\s*([\w\s]*)
```

This part of the pattern matches the name:

- `\s*` matches one or more spaces (including new lines).
- `([\w\s]*)` captures as much of the text as we can with only letters and spaces. The parentheses are used to capture what is matched inside the parenthesis. Since the name in the text is followed by the character '-', this will do.

Looping over all attributes, we append to the pattern strings such as

```
.*?(\d*\s*Crossing)
```

The `.*?` matches anything in a non-greedy fashion. The `.` stands for any character, the `*` for one or more, and the `?` does the matching in a non-greedy way, that is, it will match as little of the text as possible, as long as what follows matches `\d*\s*Crossing`. For this latter expression, the parenthesis serves to capture what is inside it. Inside it, we match any number of digits followed by any number of spaces and the attribute `'Crossing'`. We loop over all attributes to construct the full pattern.

```
a = pat.match (text)
row.append (a.group (1))
for i in range(2,len(attributes)+2):
    row.append (int (a.group(i).split()[0]))
rows.append (row)
print row[1]
```

Here we apply the pattern to the text. The returned object has stored the captured text matched inside the parentheses of our pattern, and we can access it through the `group` functions. For Messi, `'group (1)'` will have 'Lionel Messi', `'group (2)'` will have '77 Crossing', `'group (3)'` will have '95 Finishing' and so on. We use `split` to keep only the numbers.

```
df = pd.DataFrame (rows, columns = [ 'link', 'name' ] + attributes )
df.to_csv ( 'ArgentinaPlayers.csv', index = False )
```

Here we store the data in a Pandas `DataFrame` and save it as a CSV file.

3. *How would you change the code to download the first 500 English players instead?*

To download the first 500 English players instead, we need to change the code block

```
for offset in [ '0', '100', '200' ]:
    page = requests.get ( 'http://sofifa.com/players?na=52&offset=' + offset )
```

to

```
for offset in [ '0', '100', '200', '300', '400' ]:
    page = requests.get ( 'http://sofifa.com/players?na=14&offset=' + offset )
```

4. *Use the `sklearn.cluster.KMeans` Python class to cluster the players into 5 clusters.*

This can be achieved with the following code:

```
from sklearn.cluster import KMeans
import numpy as np
X = np.array (df[attributes])
kmeans = KMeans (n_clusters = 5, random_state = 0)
kmeans.fit (X)
df['label'] = pd.Series (kmeans.labels_, index = df.index)
```

5. *By inspecting the clusters and looking up individual players online, try to assign meaningful labels to the clusters.*

To print the names of the players in cluster 0, we type

```
df[df.label == 0].name
```

The first few players in cluster 0 are

1	Gonzalo Higuain
9	Mauro Icardi
45	Lucas Alario
55	Marco Ruben
59	Nicolas Blandi
61	Dario Benedetto
62	Gustavo Bou
78	Mauro Boselli
95	Franco Di Santo
100	Jonathan Calleri

116	Silvio Romero
118	Emiliano Sala
124	Maximiliano Lopez
131	German Denis
140	Facundo Ferreyra
143	Sebastian Driussi
153	Julio Furch
157	Lucas Viatri
158	Guido Carrillo
165	Jose Sand
168	Enrique Triverio
171	Juan Ignacio Gomez
172	Leonardo Ulloa
...	

A web search reveals that most of these players are strikers, so we assign to the cluster 0 the label 'Strikers'. Cluster 1 consists of the players

12	Geronimo Rulli
31	Sergio Romero
36	Marcelo Barovero
43	Nahuel Guzman
44	Willy Caballero
53	Sebastian Torrico
54	Agustin Marchesin
64	Mariano Andujar
66	Franco Armani
83	Agustin Orion
105	Mariano Barbosa
108	Fernando Monetti
123	German Lux
126	Albano Bizzarri
149	Juan Pablo Carrizo
164	Cristian Campestrini
175	Guillermo Sara
184	Luciano Pocrnjic
193	Rodrigo Rey
204	Marcos Diaz
215	Javier Garcia
229	Jorge Broun
234	Oscar Ustari
242	Julian Speroni
253	Luis Ardente
291	Nereo Fernandez

who (according to a web search) are all goalkeepers. Cluster 2 starts as follows:

7	Ever Banega
8	Javier Mascherano
11	Ezequiel Garay
13	Marcos Rojo
14	Mateo Musacchio
16	Pablo Zabaleta

17	Lucas Biglia
18	Augusto Fernandez
19	Roberto Pereyra
22	Claudio Yacob
27	Cristian Ansaldi
35	Nicolas Pareja
47	Guido Pizarro
48	Marcos Acuna
49	Enzo Perez
57	David Abraham
65	Pablo Perez
67	Lucas Castro
69	Leandro Paredes
75	Esteban Cambiasso
76	Gino Peruzzi
77	Fernando Gago
81	Facundo Roncaglia
82	Emmanuel Mas
84	Matias Kranevitter
87	Oscar Trejo
89	Ramiro Funes Mori
....	

These players are mostly defensive midfielders and full backs. Cluster 3 starts as follows:

6	Nicolas Otamendi
15	Gonzalo Rodriguez
24	Federico Fazio
25	Gustavo Cabral
28	Federico Fernandez
29	Lisandro Lopez
41	Victor Cuesta
46	Martin Demichelis
52	Santiago Gentiletti
63	Nicolas Burdisso
68	German Pezzella
72	Mauro Dos Santos
88	Jonatan Maidana
93	Santiago Vergini
106	Matias Caruzzo
109	Luciano Lollo
111	Martin Mantovani
113	Julio Alberto Barroso
120	Jonathan Schunke
125	Nicolas Spolli
137	Marcos Angeleri
139	Renato Civelli
145	Juan Insaurralde
148	Carlos Izquierdoz
152	Matias Zaldivia
159	Jose Maria Basanta
166	Leandro Desabato

```

174         Juan Forlin
177         Fernando Tobio
...

```

These players are almost exclusively central defenders. Finally, cluster 4 starts as follows:

```

0         Lionel Messi
2         Sergio Aguero
3         Angel Di Maria
4         Paulo Dybala
5         Nicolas Gaitan
10        Javier Pastore
20        Erik Lamela
21        Alejandro Gomez
23        Diego Perotti
26        Fernando Belluschi
30        Manuel Lanzini
32        Luciano Vietto
33        Pablo Batalla
34        Lisandro Lopez
37        Rodrigo Palacio
38        Angel Correa
39        Ignacio Piatti
40        Eduardo Salvio
42        Jose Sosa
50        Diego Valeri
51        Pablo Piatti
56        Diego Buonanotte
58        Rogelio Funes Mori
60        Pablo De Blasis
70        Franco Cervi
71        Mauro Zarate
73        Lautaro Acosta
74        Sebastian Blanco
...

```

These players are mostly playmakers/attacking midfielders, or forwards that like to move out of the box. To summarize, we obtain the following clusters:

Cluster	Description
0	Strikers
1	Goalkeepers
2	Defensive midfielders and full backs
3	Central defenders
4	Playmakers/attacking midfielders, outside forwards

It is interesting that KMeans has found side backs to be closer related to defensive midfielders than central defenders. Also, there was a clear distinction between strikers and other types of forwards that were grouped together with attacking midfielders.

6. For a new and unknown player, the following attributes are available: (...) For each of your 5 clusters from Step 4, compute the cluster centroid. Assign the new player to the

nearest cluster based on the distance to the cluster centroids, using only the available attributes.

To decide which cluster to assign the new player to, we execute the following code:

```
centers = kmeans.cluster_centers_
attributesofinterest = [ 'Crossing', 'Sprint_Speed', 'Long_Shots', 'Aggression',
    , 'Marking', 'Finishing', 'GK_Handling' ]
indices = [i for i in range (len (attributes)) if attributes[i] in
    attributesofinterest]
centersofinterest = centers[:, indices]
playerattributes = np.array ([45, 40, 35, 45, 60, 40, 15])
for cluster in range (5):
    dis = np.sqrt (((centersofinterest[cluster, :] - playerattributes) ** 2).
        sum())
    print cluster, dis
```

We obtain the following result:

```
0 59.2088045066
1 86.925706223
2 57.3034127353
3 46.3560352058
4 64.7386383339
```

Our best guess is therefore that the player is a central defender.